

Introduction to Stateful Stream Processing with Apache Flink

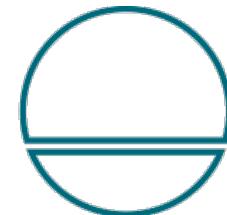


Robert Metzger
@rmetzger_
robert@ververica.com





Original creators of
Apache Flink®



Ververica Platform
Open Source Apache Flink
+ Application Manager



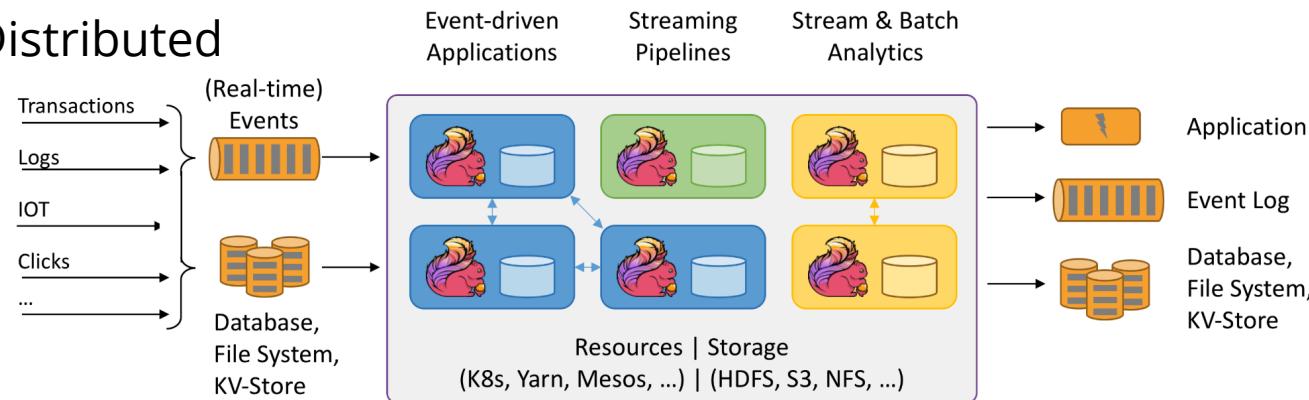
Apache Flink 101



Apache Flink



- Apache Flink is an open source stream processing framework
 - Low latency
 - High throughput
 - Stateful
 - Distributed



What is Apache Flink?



Batch Processing

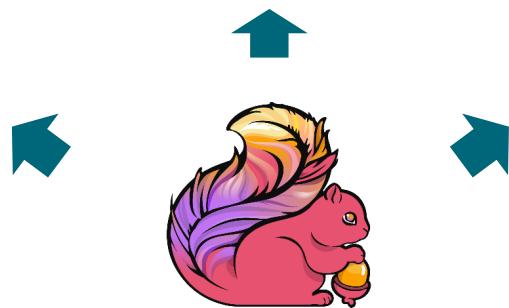
process static and historic data

Data Stream Processing

realtime results from data streams

Event-driven Applications

data-driven actions and services



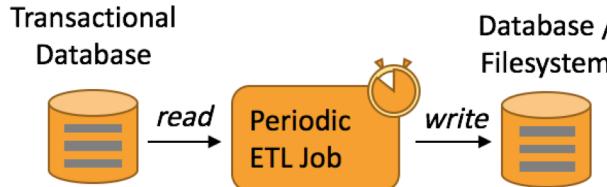
Stateful Computations Over Data Streams



Use Case: Streaming ETL

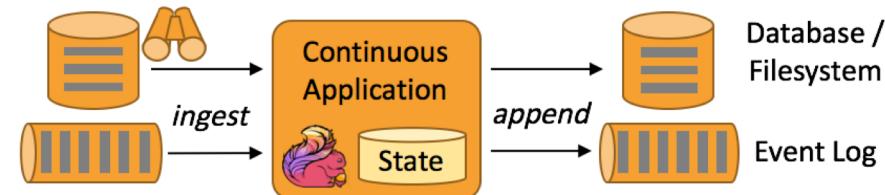
- Periodic ETL is the traditional approach
 - External tool periodically triggers ETL batch job
- Data pipelines continuously move data
 - Ingestion with low latency
 - No artificial data boundaries

Periodic ETL



Data Pipeline / Real-time ETL

Real-time Events

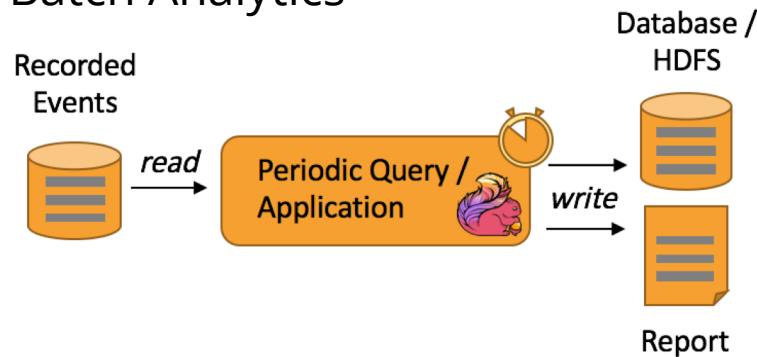




Use Case: Data Analytics

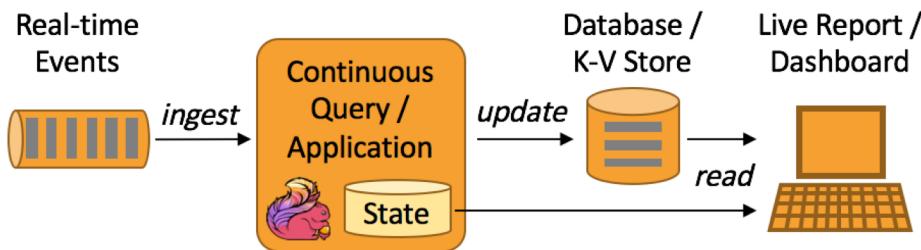
- Batch analytics is great for ad-hoc queries
 - Queries change faster than data
 - Interactive analytics / prototyping

Batch Analytics



- Stream analytics continuously processes data
 - Data changes faster than queries
 - Live / low latency results
 - No Lambda architecture required!

Stream Analytics

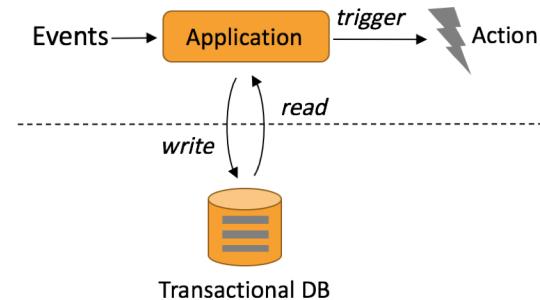


Use Case: Event-driven applications

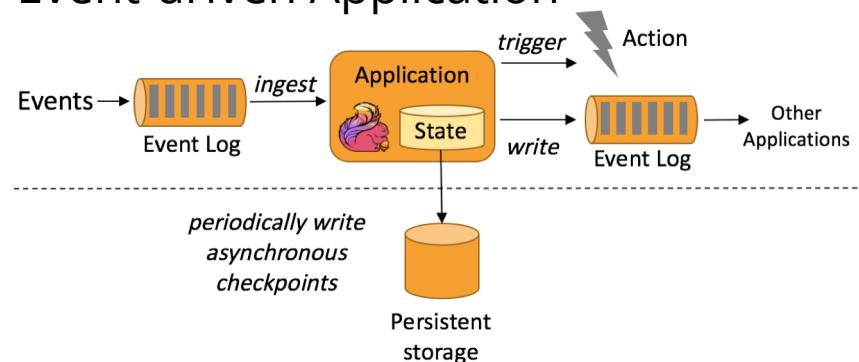


- Traditional application design
 - Compute & data tier architecture
 - React to and process events
 - State is stored in (remote) database
- Event-driven application
 - State is maintained locally
 - Guaranteed consistency by periodic state checkpoints
 - Tight coupling of logic and data (microservice architecture)
 - Highly scalable design

Transactional Application



Event-driven Application



Hardened at scale



Details about their use cases and more users are listed on Flink's website at <https://flink.apache.org/poweredsby.html>

Case Study: Single's Day



- Chinese Shopping Festival
- Very high peak load
 - 100s millions records per second
 - 100s thousands payments per second
 - 10 TBs of managed state
 - 10s thousands of cores
- Flink used in various areas in the process incl. payment, shipping, realtime recommendations and the giant dashboard



References

- <https://www.ververica.com/blog/singles-day-2018-data-in-a-flink-of-an-eye>
- https://medium.com/@alitech_2017/how-to-cope-with-peak-data-traffic-on-11-11-the-secrets-of-alibaba-stream-computing-17d5e807980c



Building blocks



The Core Building Blocks

Event Streams

real-time and hindsight

State

complex business logic

(Event) Time

consistency with out-of-order data and late data

Snapshots

forking / versioning / time-travel

Stateful Event & Stream Processing



```
val lines: DataStream[String] = env.addSource(new FlinkKafkaConsumer(...))
```

} *Source*

```
val events: DataStream[Event] = lines.map((line) => parse(line))
```

} *Transformation*

```
val stats: DataStream[Statistic] = stream  
  .keyBy("sensor")  
  .timeWindow(Time.seconds(5))  
  .sum(new MyAggregationFunction())
```

} *Transformation*

```
stats.addSink(new RollingSink(path))
```

} *Sink*



Source



Transform



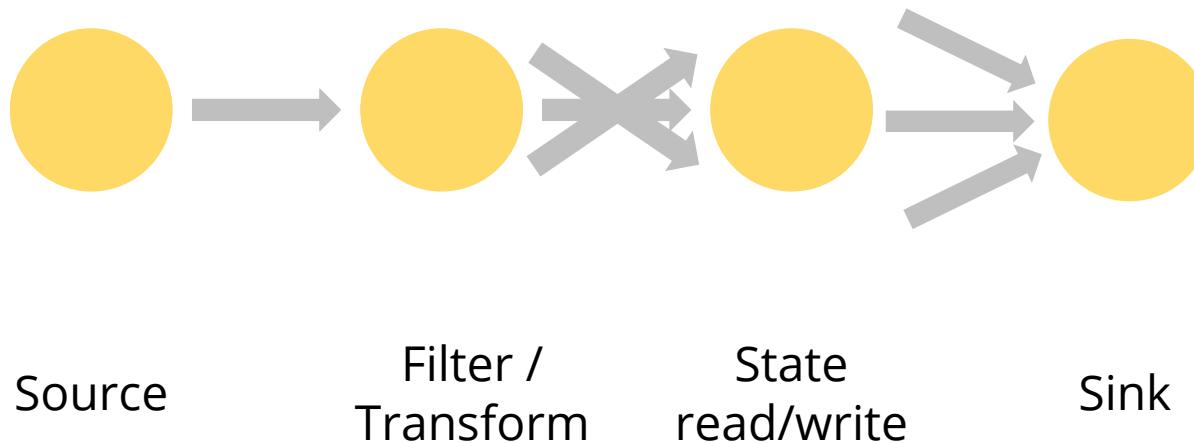
Window
(state read/write)



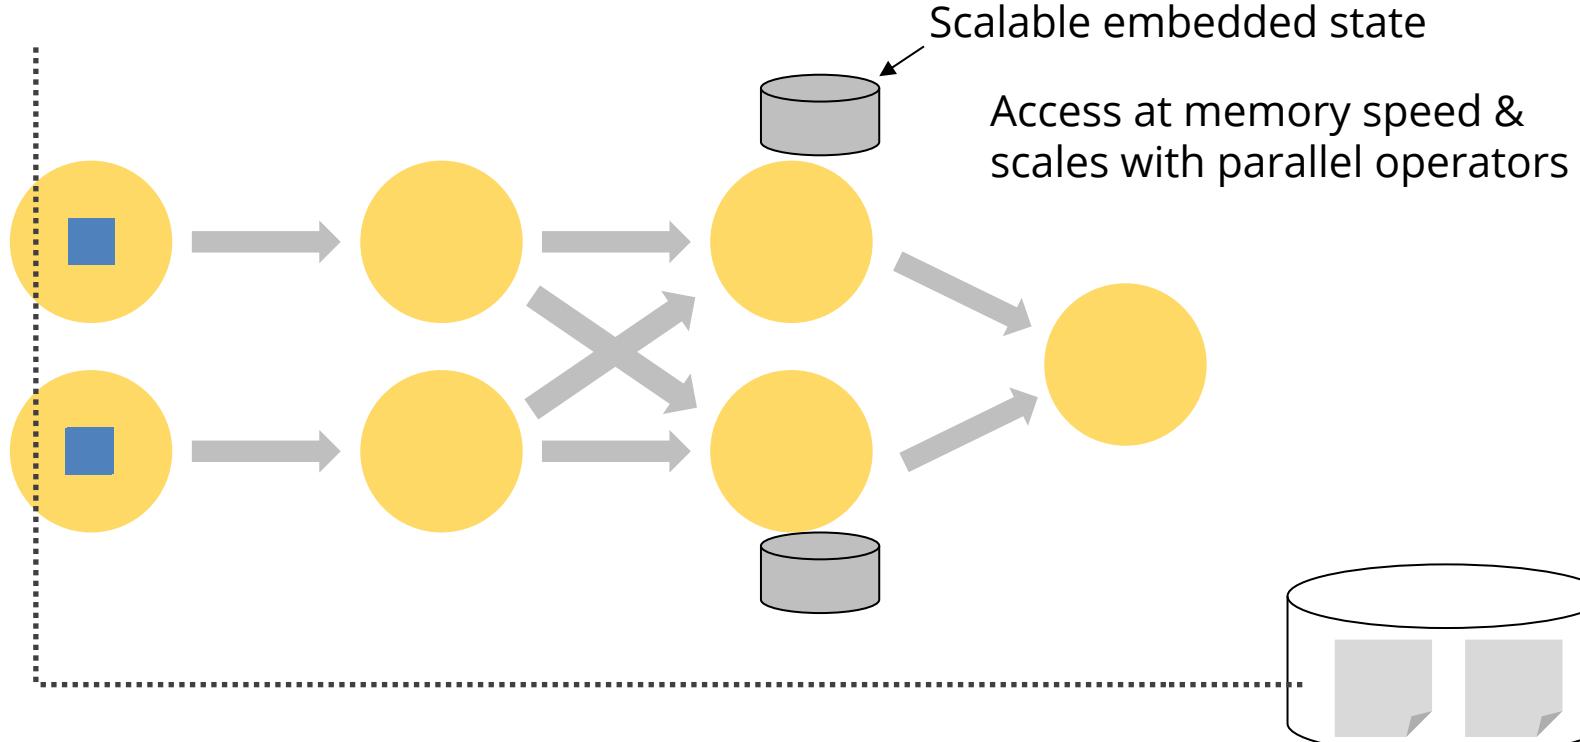
Sink

} *Streaming Dataflow*

Stateful Event & Stream Processing



Stateful Event & Stream Processing

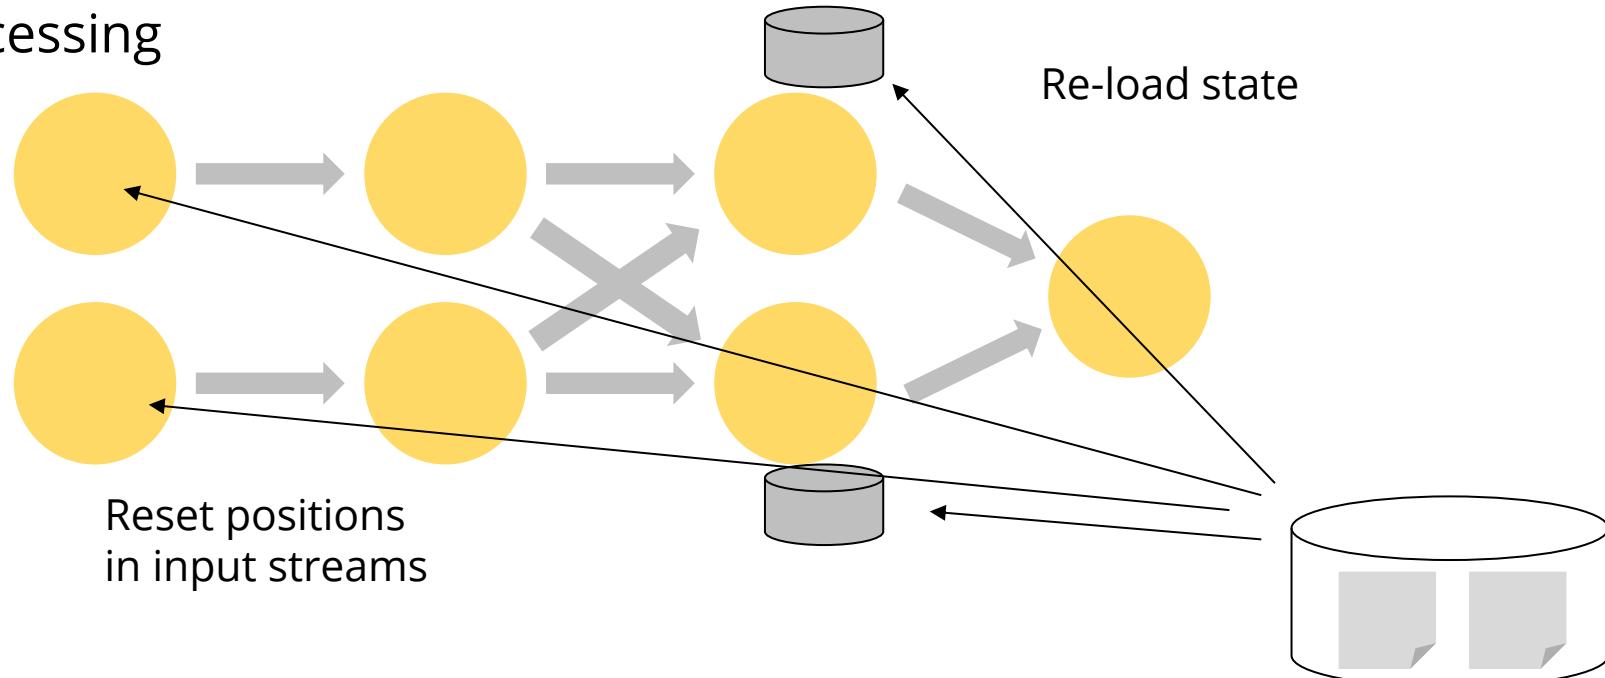


Stateful Event & Stream Processing

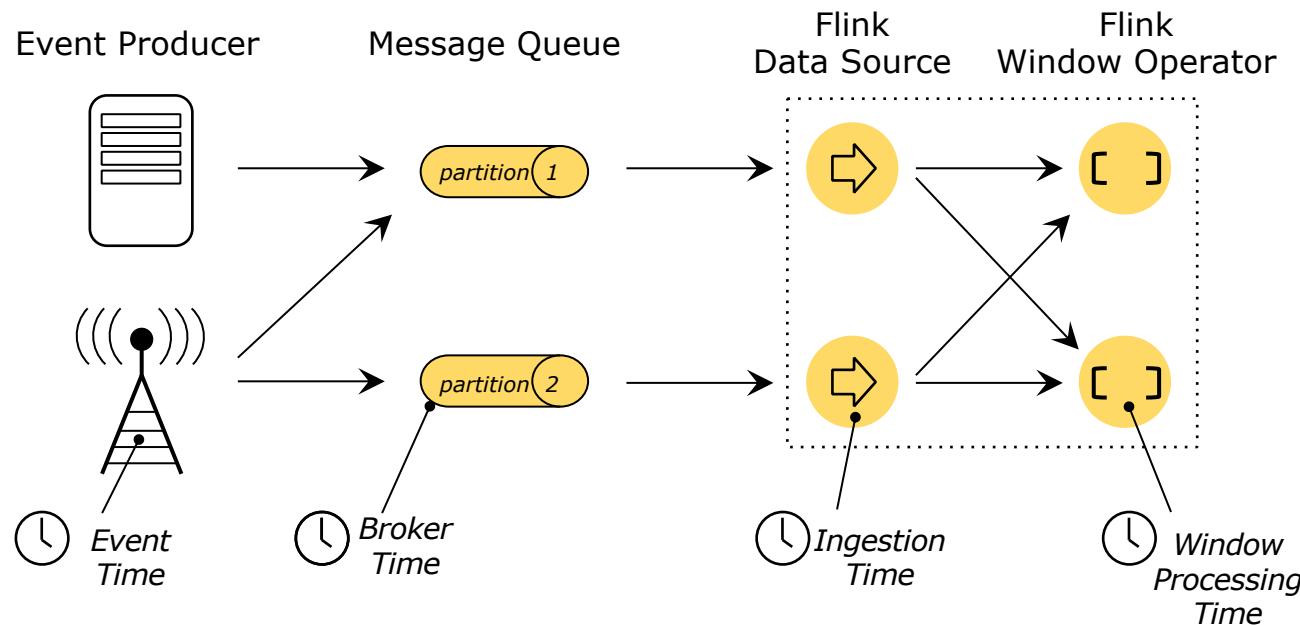


Rolling back computation

Re-processing



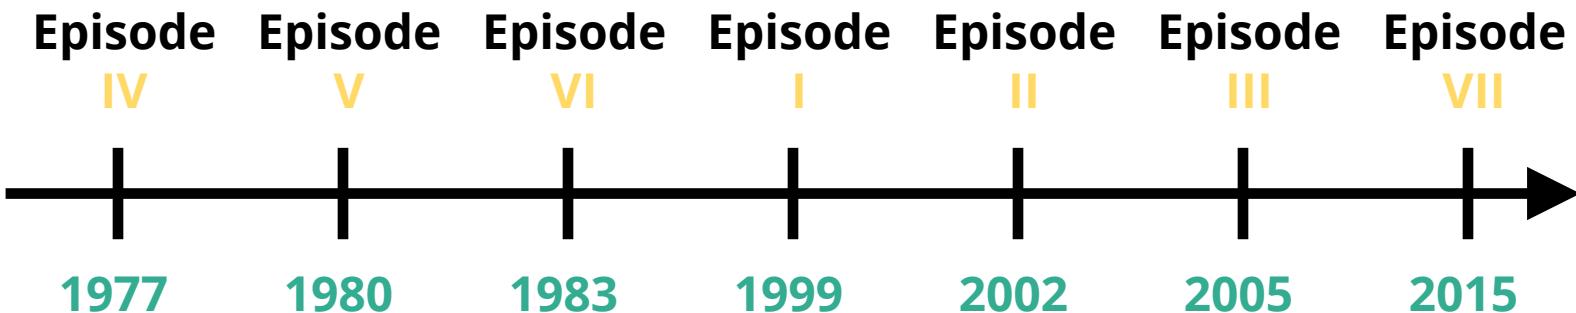
Time: Different Notions of Time



Time: Event Time Example



Event Time



Processing Time



Recap: The Core Building Blocks

Event Streams

real-time and hindsight

State

complex business logic

(Event) Time

consistency with out-of-order data and late data

Snapshots

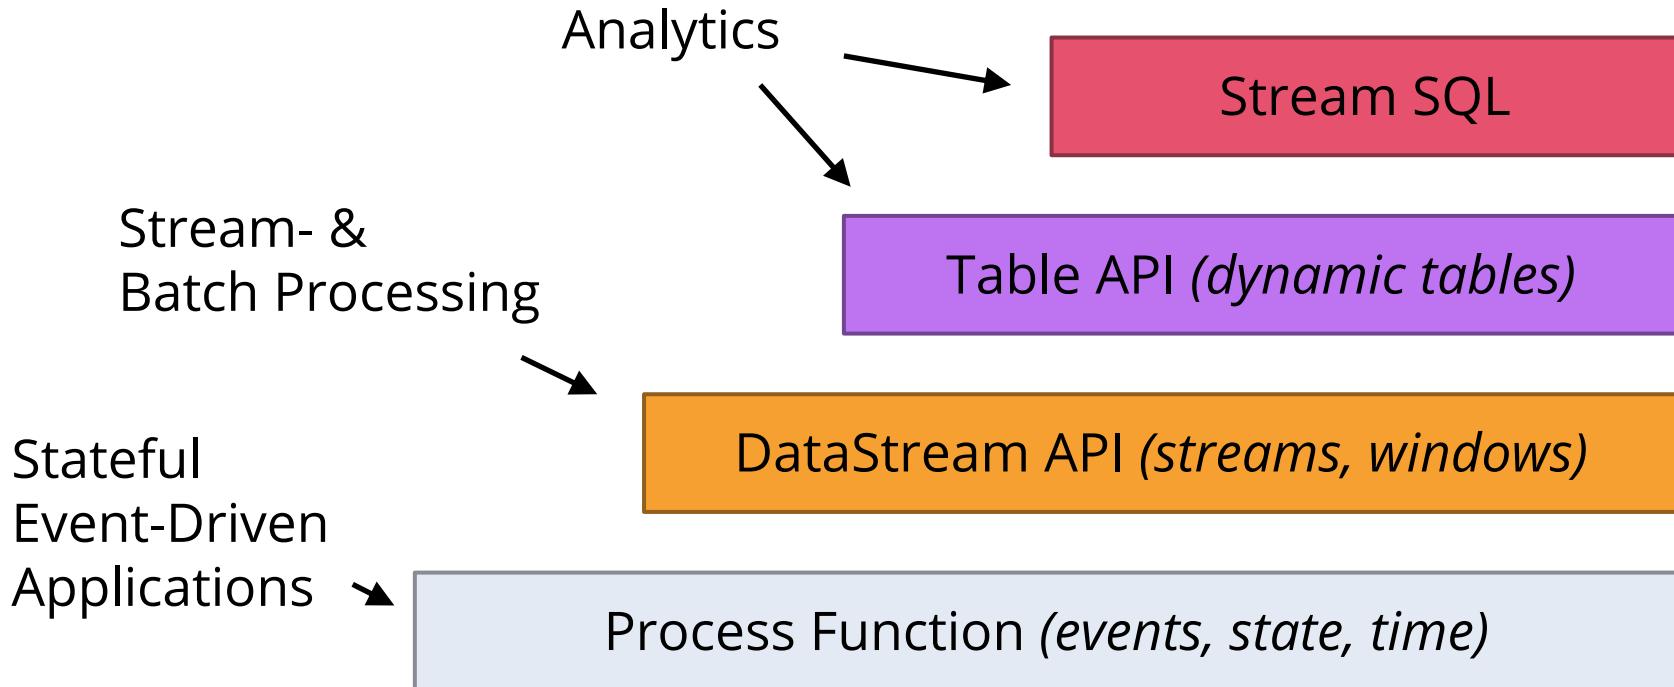
forking / versioning / time-travel



APIs



The APIs



Process Function



```
class MyFunction extends ProcessFunction[MyEvent, Result] {

    // declare state to use in the program
    lazy val state: ValueState[CountWithTimestamp] = getRuntimeContext().getState(...)

    def processElement(event: MyEvent, ctx: Context, out: Collector[Result]): Unit = {
        // work with event and state
        (event, state.value) match { ... }

        out.collect(...) // emit events
        state.update(...) // modify state

        // schedule a timer callback
        ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
    }

    def onTimer(timestamp: Long, ctx: OnTimerContext, out: Collector[Result]): Unit = {
        // handle callback when event-/processing- time instant is reached
    }
}
```

Data Stream API



```
val lines: DataStream[String] = env.addSource(  
    new FlinkKafkaConsumer<>(...))  
  
val events: DataStream[Event] = lines.map((line) => parse(line))  
  
val stats: DataStream[Statistic] = stream  
    .keyBy("sensor")  
    .timeWindow(Time.seconds(5))  
    .sum(new MyAggregationFunction())  
  
stats.addSink(new RollingSink(path))
```

Table API & Stream SQL



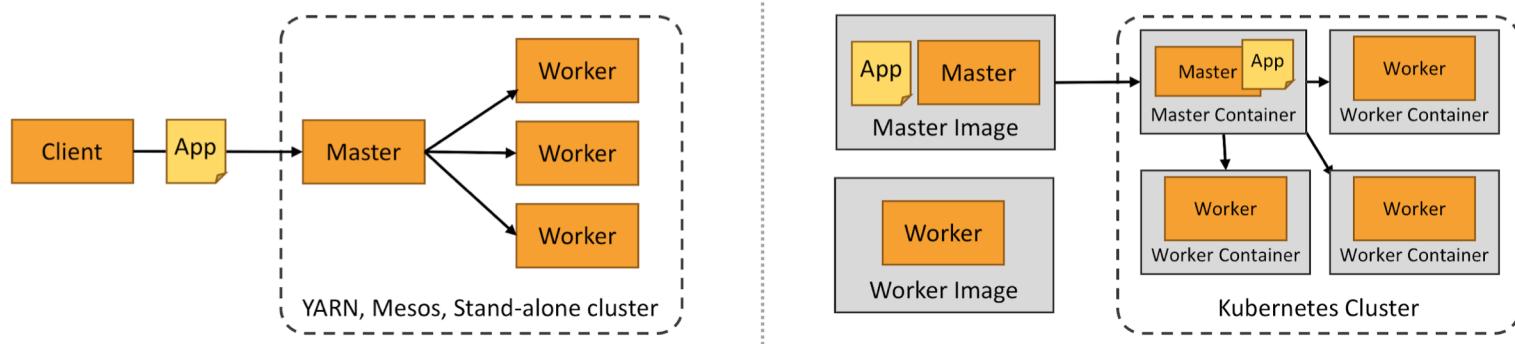
```
// Table API
val tapiResult: Table = tEnv.scan("sensors")    // scan sensors table
  .window(Tumble over 1.hour on 'rowtime' as 'w') // define 1-hour window
  .groupBy('w, 'room)                          // group by window and room
  .select('room, 'w.end, 'temp.avg as 'avgTemp) // compute average temperature
```

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp) AS avgTemp
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```



Deployment & Integrations

Deployment



kubernetes



aws



MESOS

eventador.io

Integrations



- Event logs:
 - Kafka, Kinesis, Pulsar*
- File systems:
 - S3, HDFS, NFS, MapR FS, ...
- Encodings:
 - Avro, JSON, CSV, ORC, Parquet
- Databases:
 - JDBC, HCatalog
- Key-Value Stores
 - Cassandra, Elasticsearch, Redis*





Concluding...

The Apache Flink® Conference

Berlin | October 7-9, 2019

Early Bird ticket sales ends July 15th

flink-forward.org



Organized by  ververica

#flinkforward



Q & A

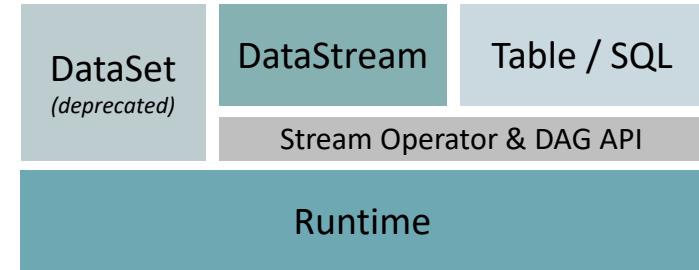
Get in touch via eMail:
robert@ververica.com
info@ververica.com

Get in touch via Twitter:
[@rmetzger_](https://twitter.com/rmetzger_)
[@ApacheFlink](https://twitter.com/ApacheFlink)

What's happening in Flink these days ...



- INSERT INTO flink_sql SELECT * FROM blink_sql
 - Turning Table API into an API unified across batch and streaming (FLINK-11439)
 - Integration with Hive ecosystem (FLINK-10556)
- Batch runtime improvements: Fine-grained recovery (FLINK-4256), more schedulers (FLINK-10429), pluggable shuffle service (FLINK-10653)
- Machine Learning Pipelines on Table API (FLIP-39)
- Table API: Caching of intermediate results (.cache() API) (FLINK-11199)
- Table API: Python support (FLINK-12308)



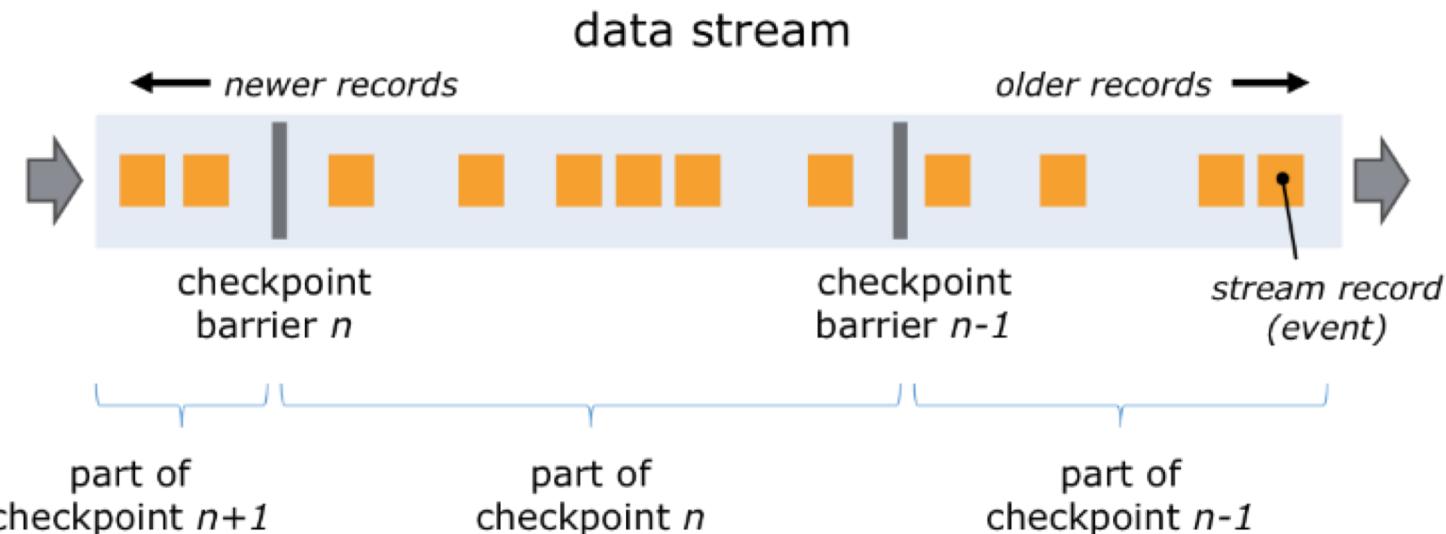


Implementation: State Checkpointing

State, Snapshots, Recovery



Coordination via markers, injected into the streams





Master

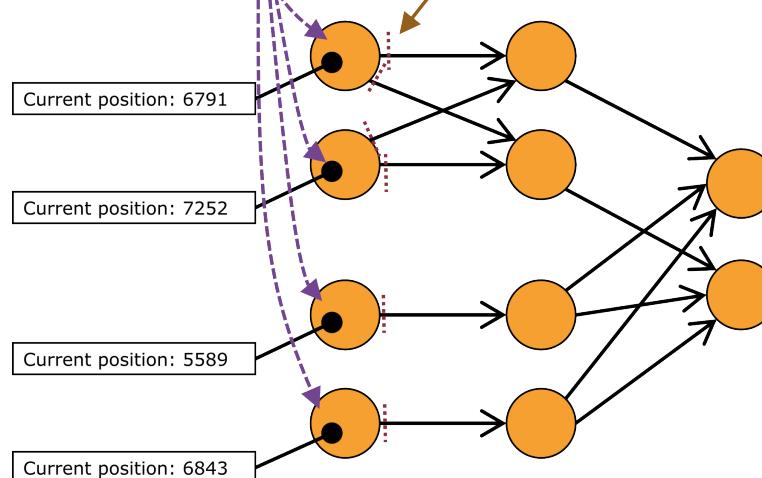
Checkpoint data

Source 1:	State 1:
Source 2:	State 2:
Source 3:	Sink 1: (pending)
Source 4:	Sink 2: (pending)

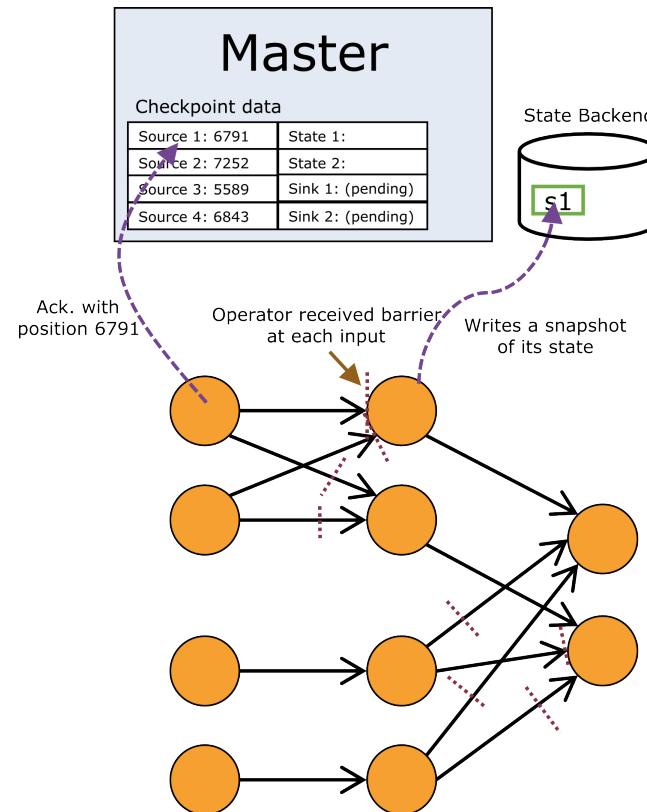
State Backend

Start checkpoint message

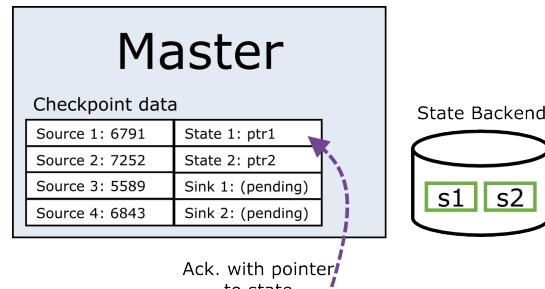
Emit stream barriers



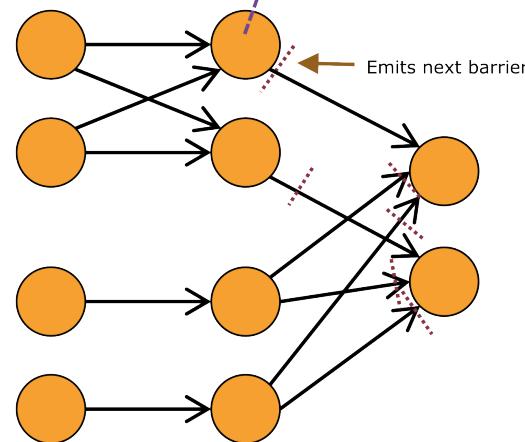
Starting
Checkpoint



Checkpoint
in Progress



Ack. with pointer
to state



Checkpoint
in Progress

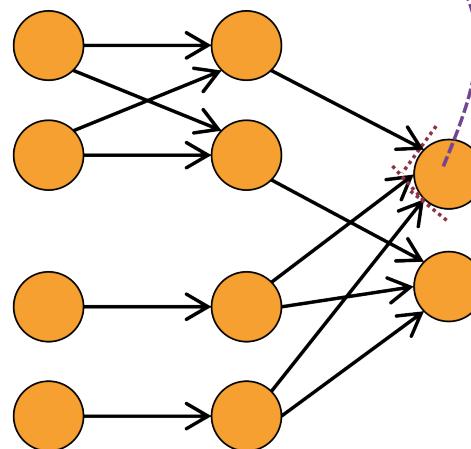


Master

Checkpoint data

Source 1: 6791	State 1: ptr1
Source 2: 7252	State 2: ptr2
Source 3: 5589	Sink 1: ack!
Source 4: 6843	Sink 2: ack!

State Backend



Sink acknowledges
checkpoint after
receiving all barriers

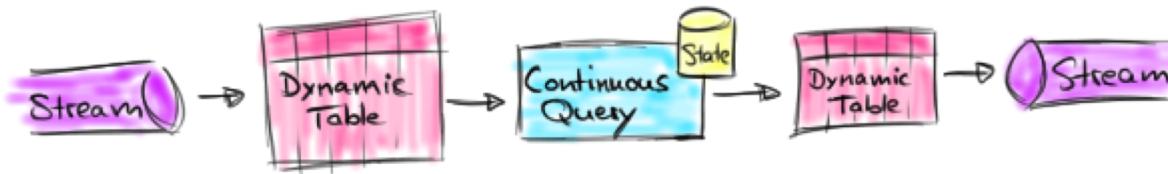
Checkpoint
Completed



Implementation: Stream SQL



Stream SQL: Intuition



Stream / Table Duality

Table without Primary Key

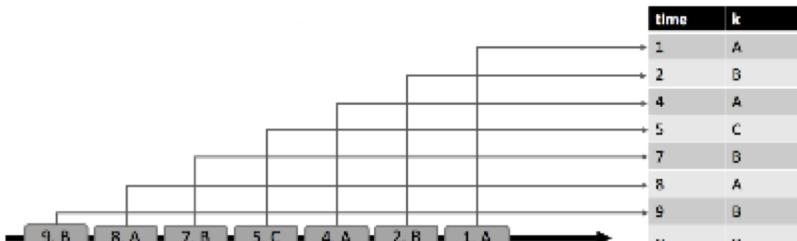
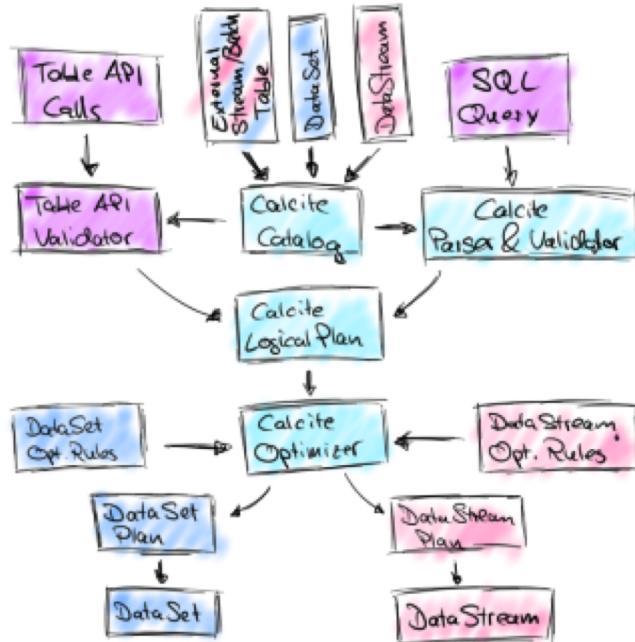


Table with Primary Key

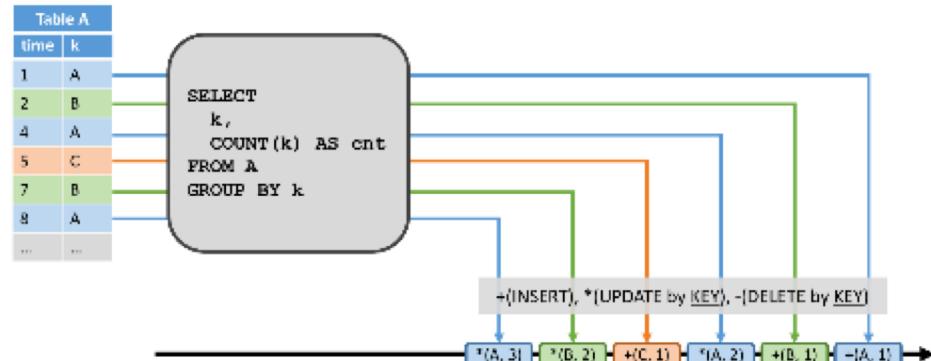




Stream SQL: Implementation



Query Compilation



Differential Computation
(add/mod/del)

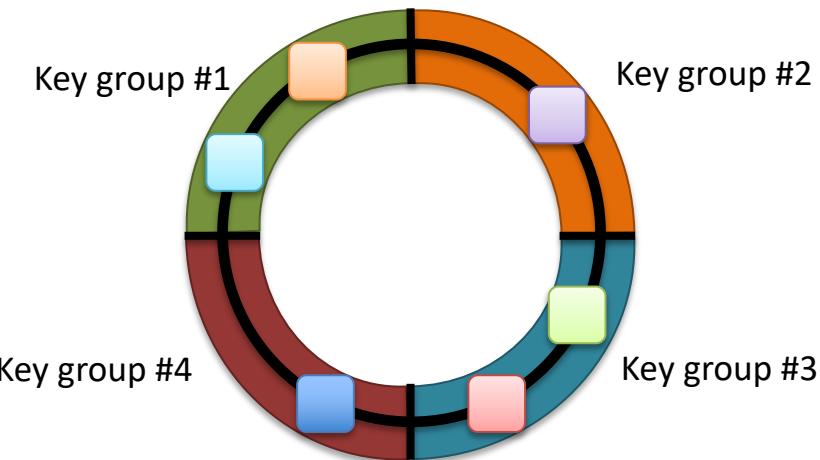


Implementation: Rescaling

Rescaling State / Elasticity



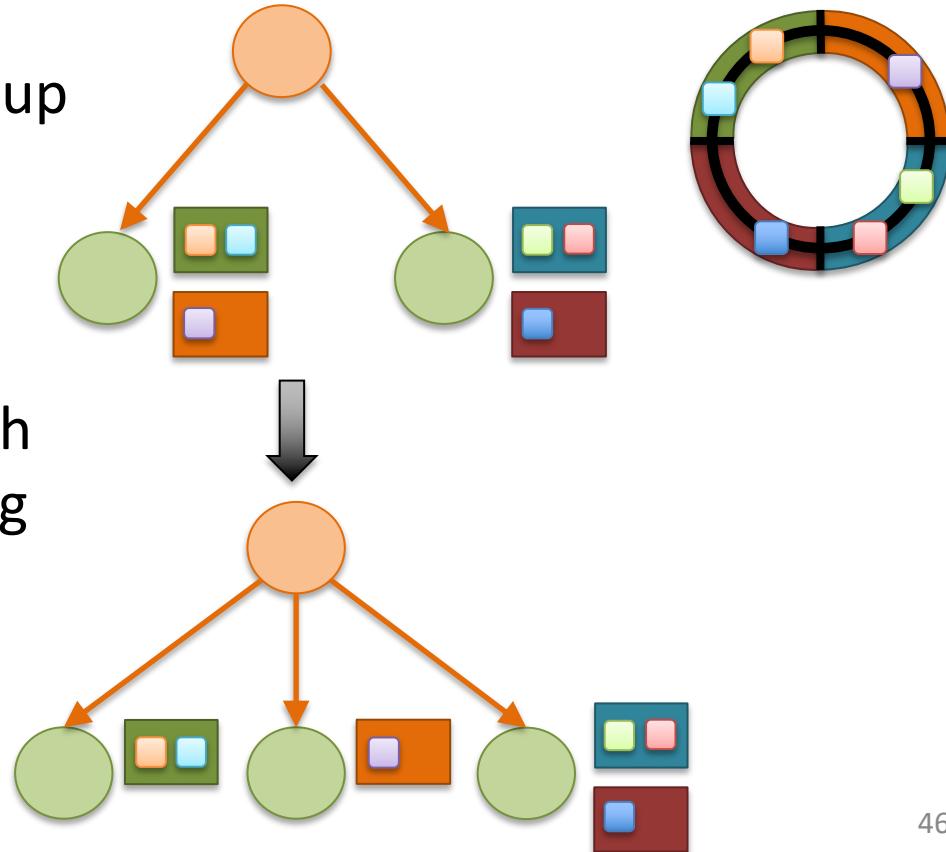
- Similar to consistent hashing Key space
- Split key space into key groups
- Assign key groups to tasks



Rescaling State / Elasticity



- Rescaling changes key group assignment
- Maximum parallelism defined by #key groups
- Rescaling happens through restoring a savepoint using the new parallelism

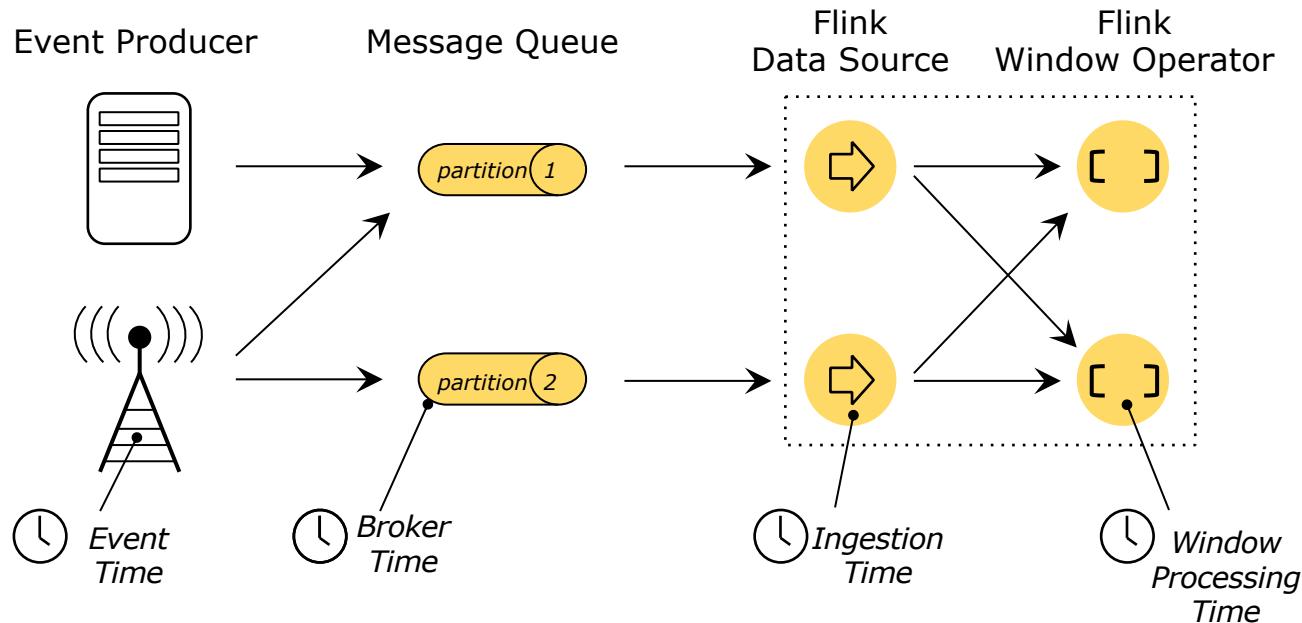




Implementation: Time-handling

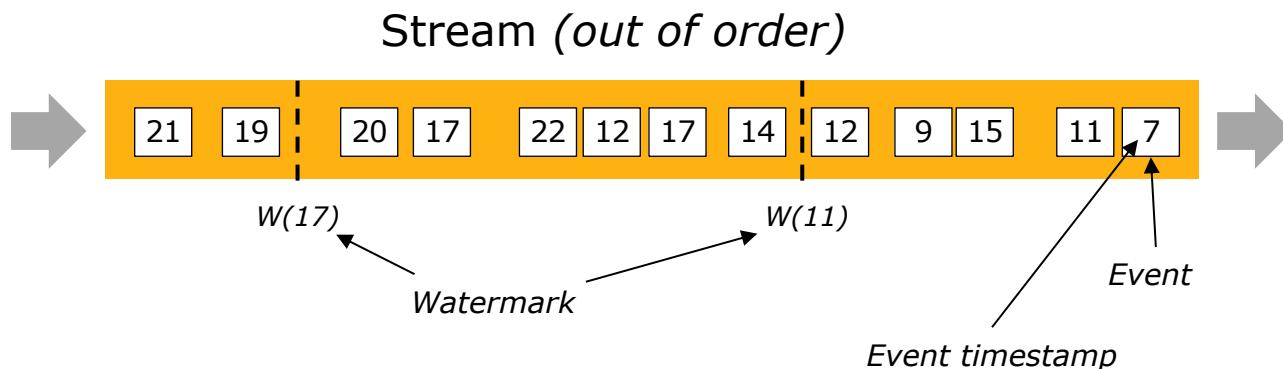
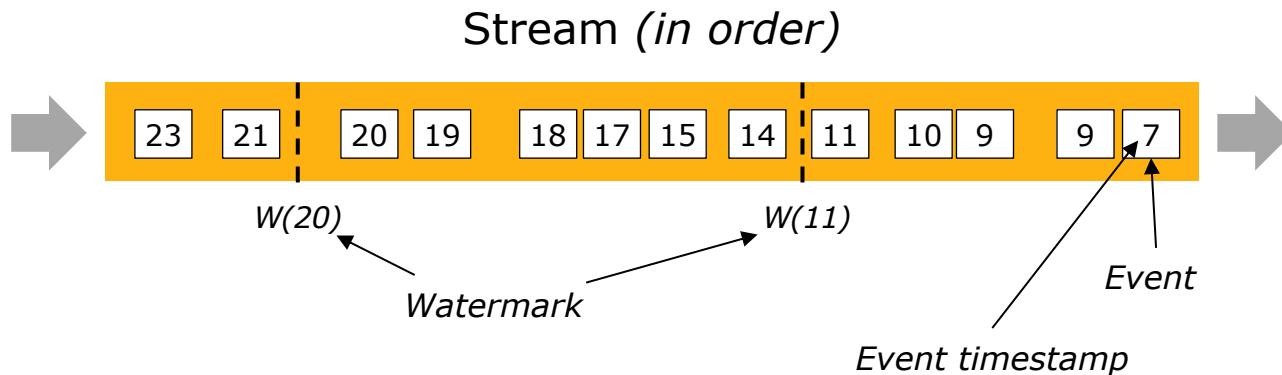


Time: Different Notions of Time

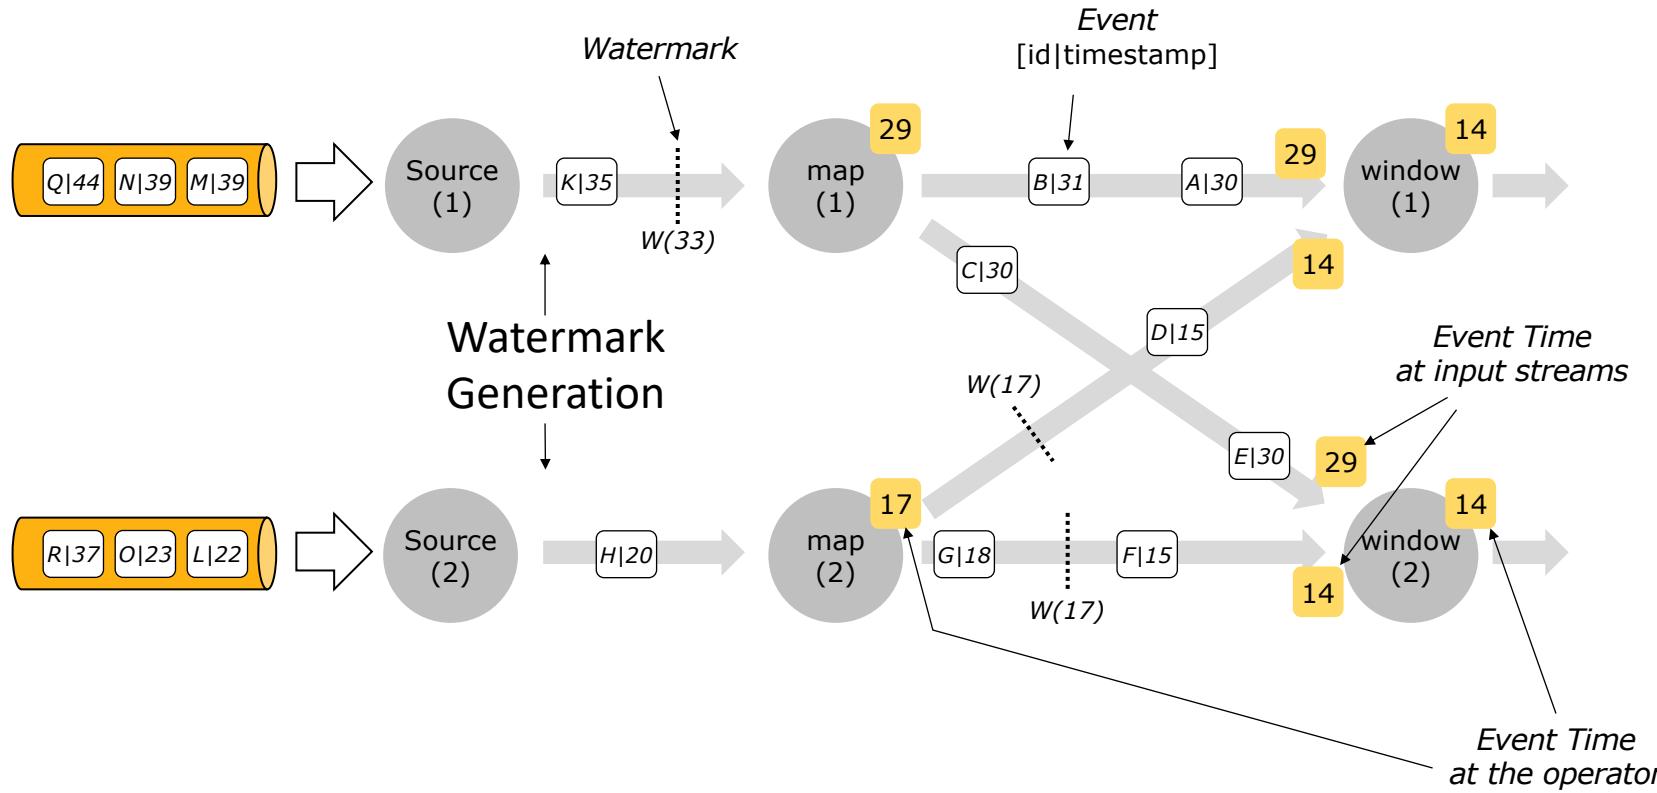




Time: Watermarks



Time: Watermarks in Parallel

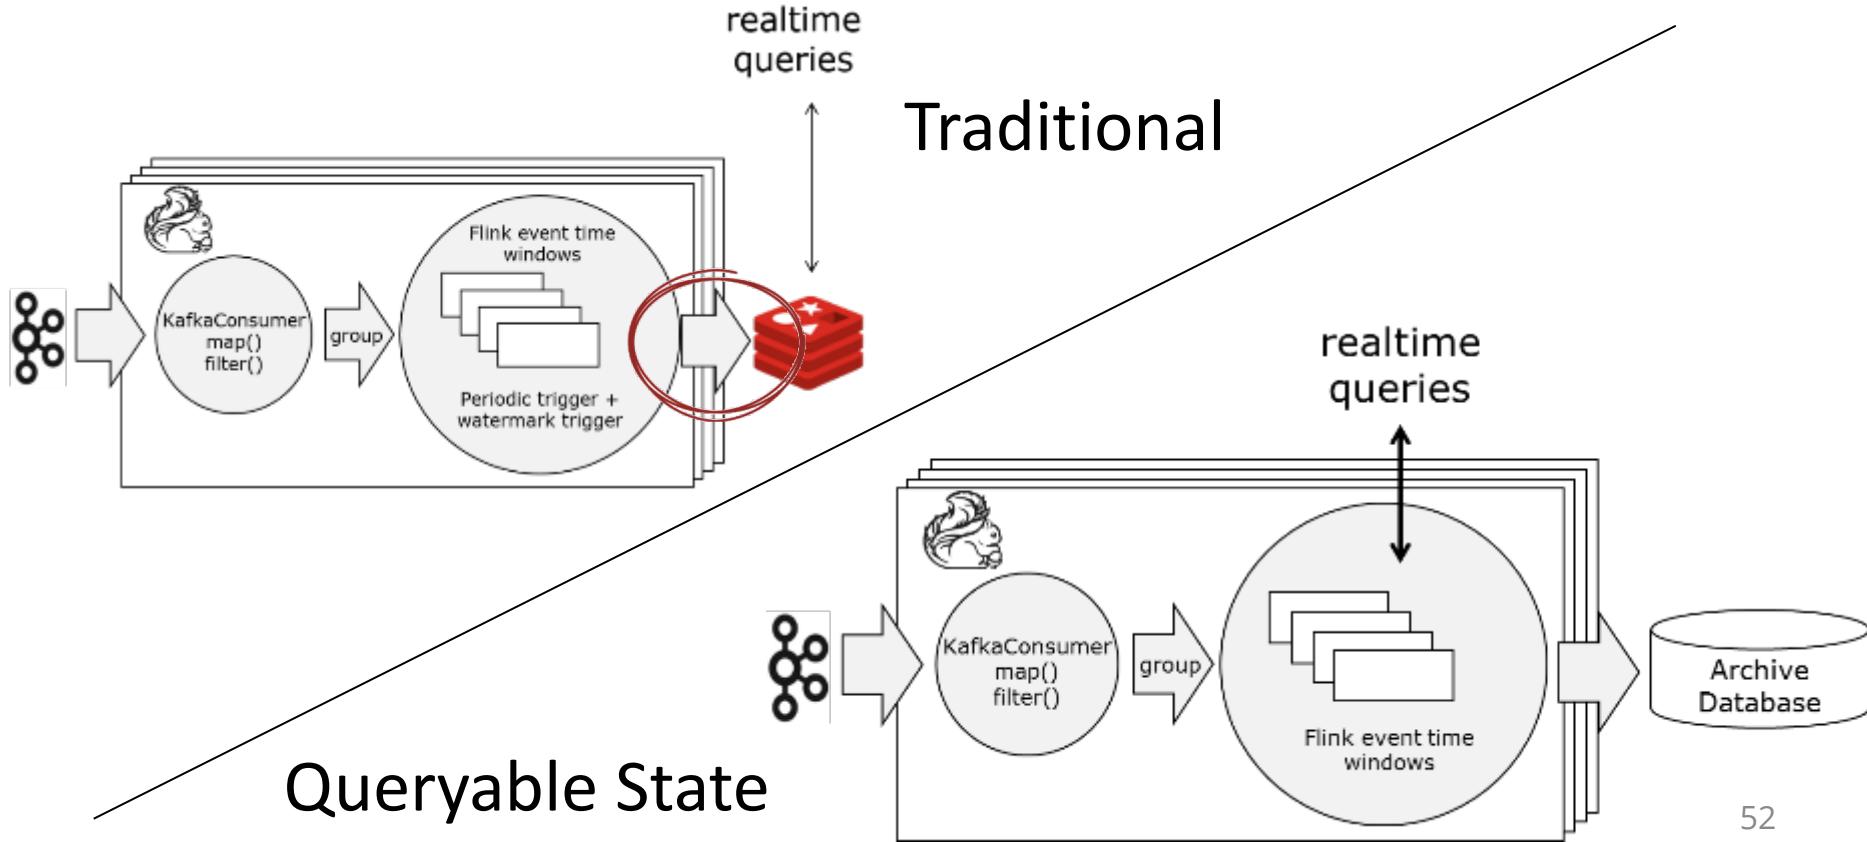




Implementation: Queryable State



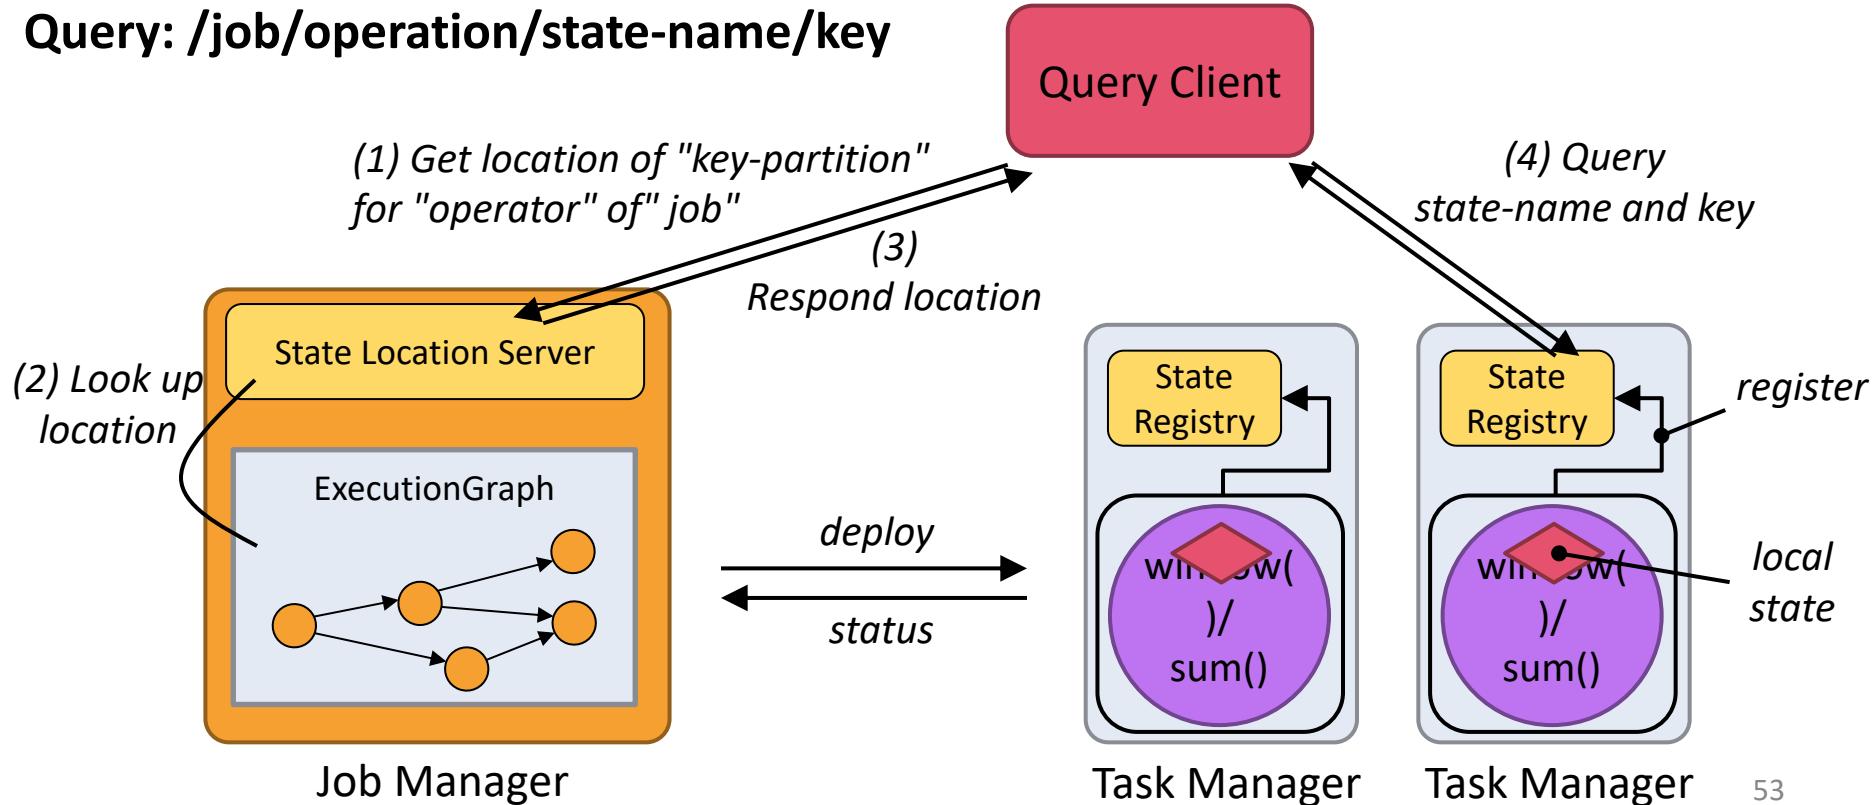
Queryable State



Queryable State: Implementation



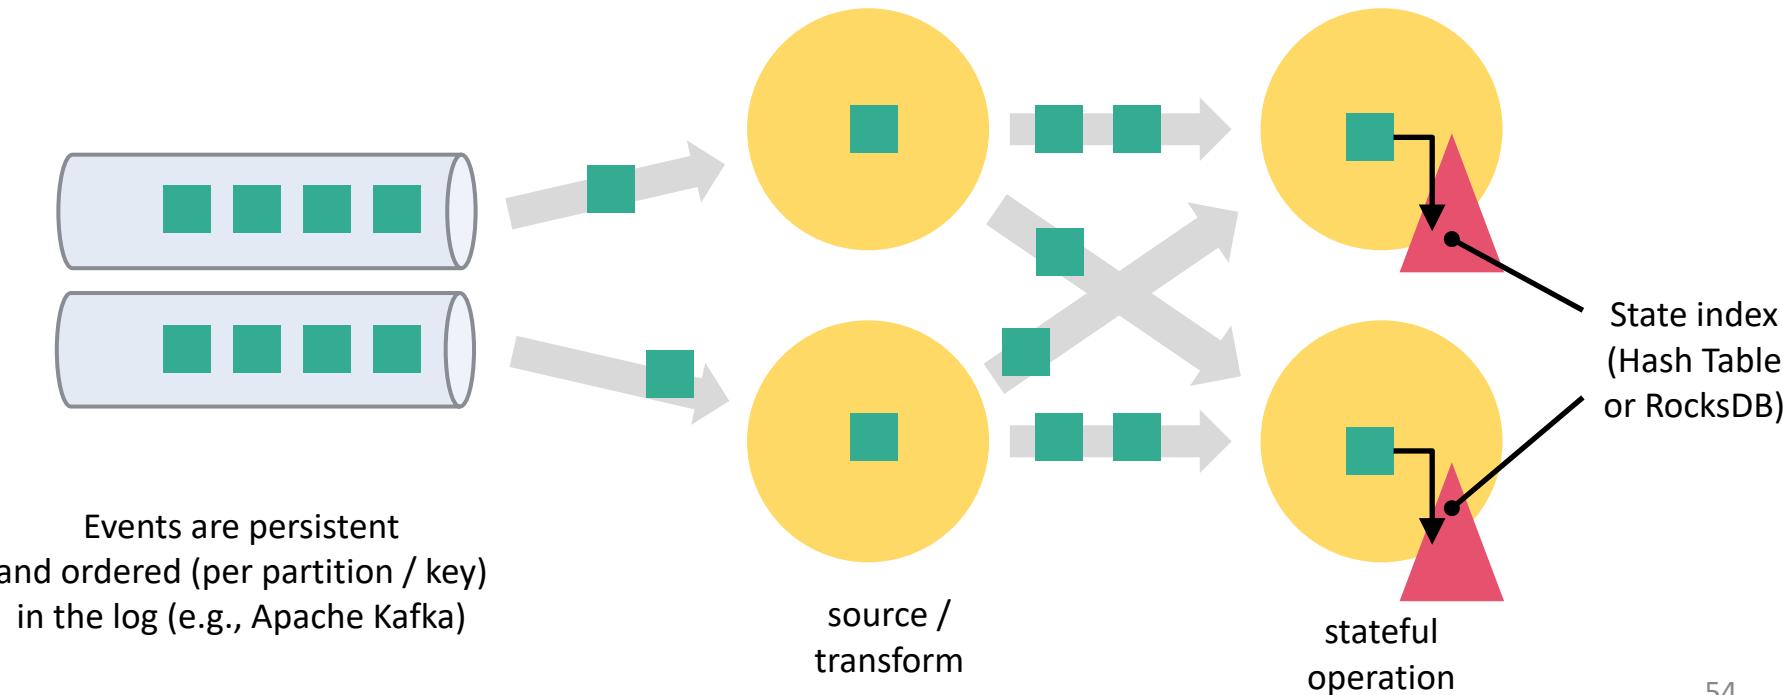
Query: /job/operation/state-name/key





State, Snapshots, Recovery

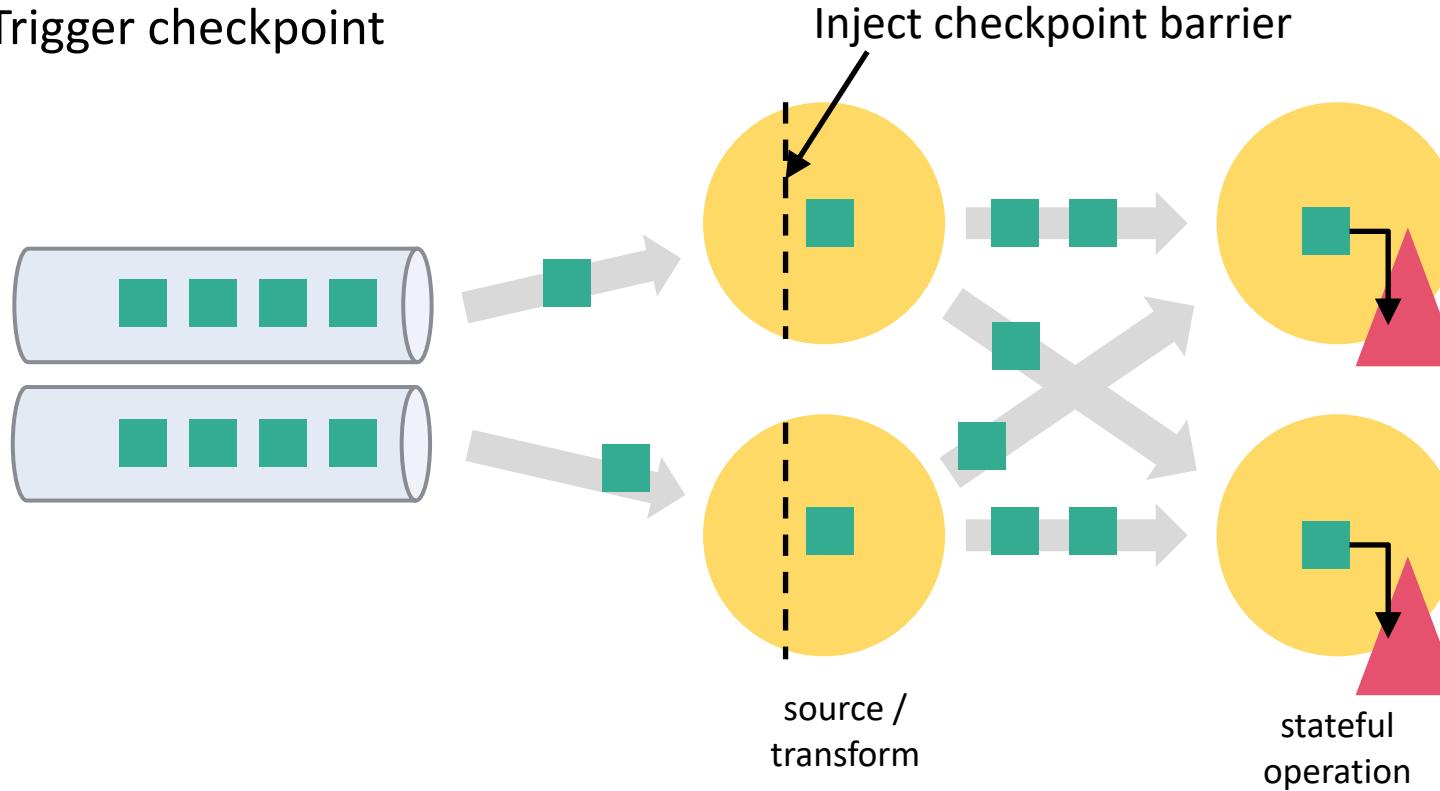
Events flow without replication or synchronous writes





State, Snapshots, Recovery

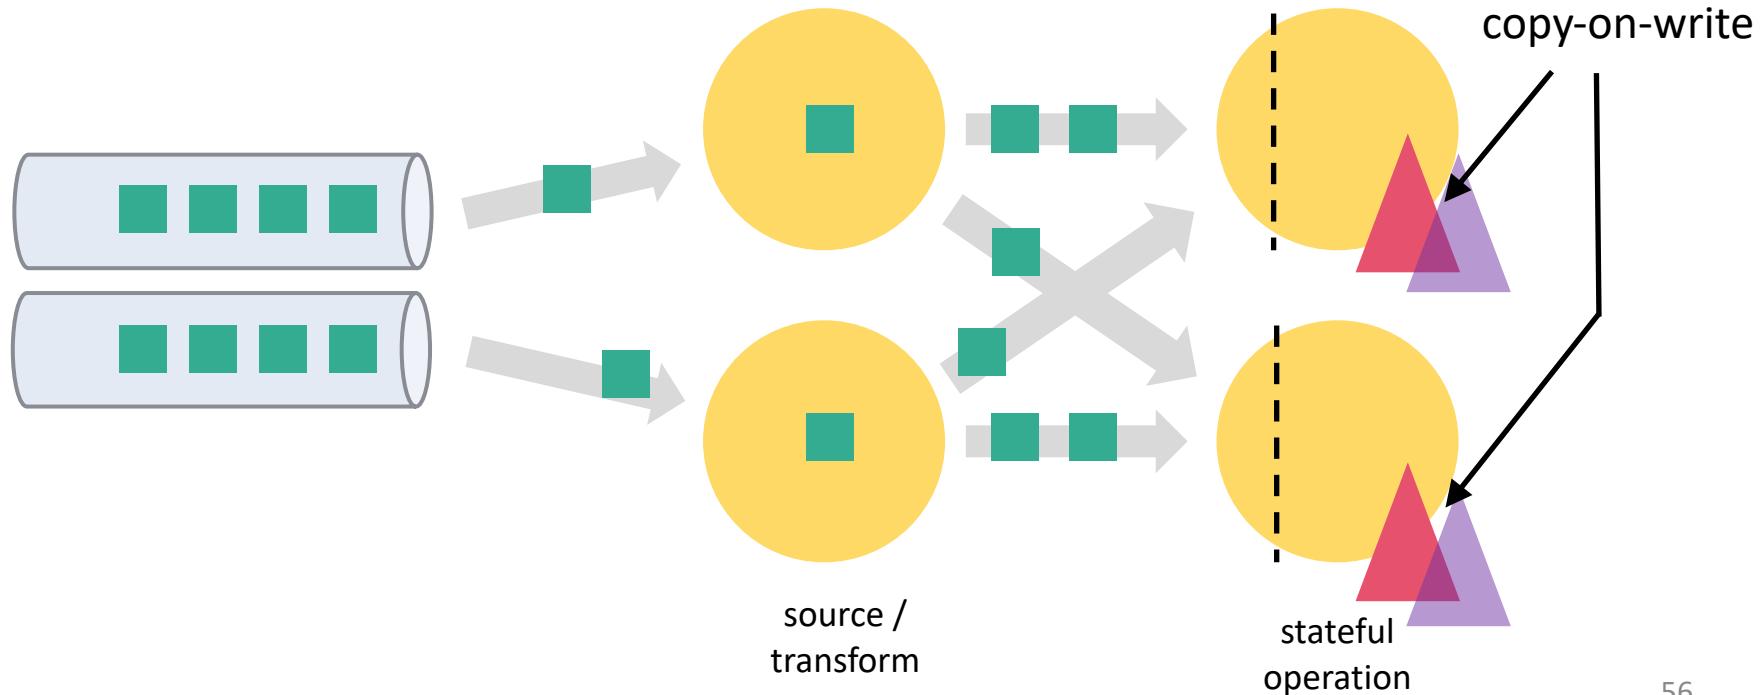
Trigger checkpoint



State, Snapshots, Recovery



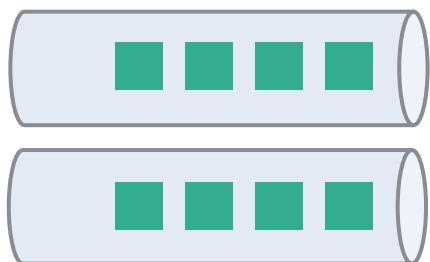
Take state snapshot



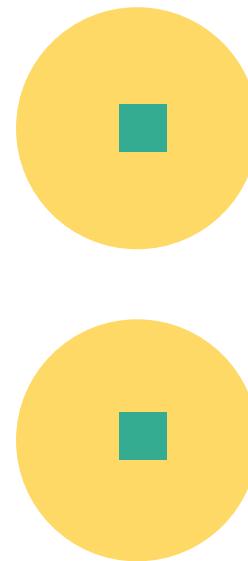
State, Snapshots, Recovery



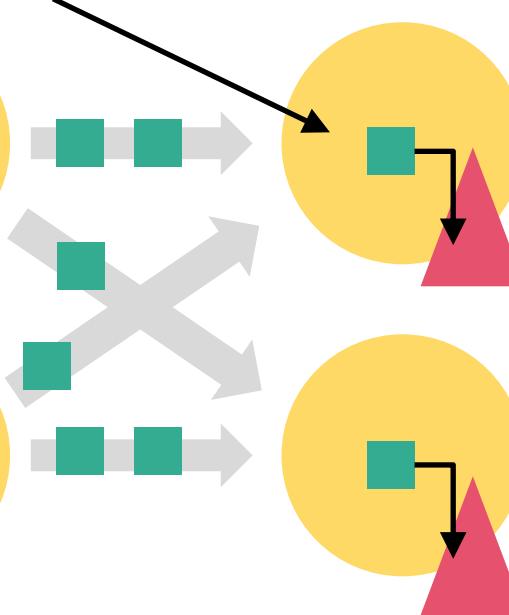
Persist state snapshots



Processing pipeline continues

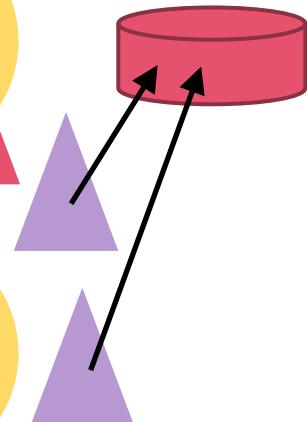


source /
transform



stateful
operation

Durably persist
snapshots
asynchronously



Powerful Abstractions



Layered abstractions to
navigate simple to complex use cases

High-level
Analytics API

Stream SQL / Tables (*dynamic tables*)

```
SELECT room, TUMBLE_END(rowtime, INTERVAL '1' HOUR), AVG(temp)
FROM sensors
GROUP BY TUMBLE(rowtime, INTERVAL '1' HOUR), room
```

Stream- & Batch
Data Processing

DataStream API (*streams, windows*)

```
val stats = stream
.keyBy("sensor")
.timeWindow(Time.seconds(5))
.sum((a, b) -> a.add(b))
```

Stateful Event-
Driven Applications

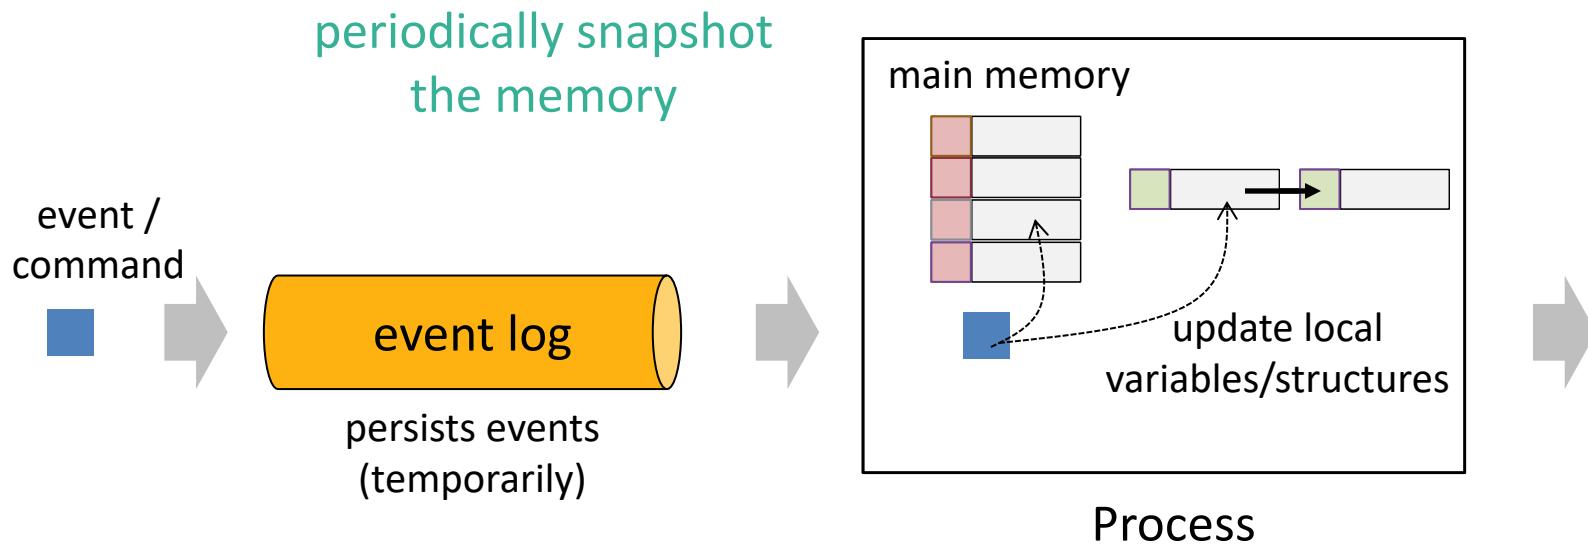
Process Function (*events, state, time*)

```
def processElement(event: MyEvent, ctx: Context, out: Collector[Result]) = {
    // work with event and state
    (event, state.value) match { ... }

    out.collect(...) // emit events
    state.update(...) // modify state

    // schedule a timer callback
    ctx.timerService.registerEventTimeTimer(event.timestamp + 500)
}
```

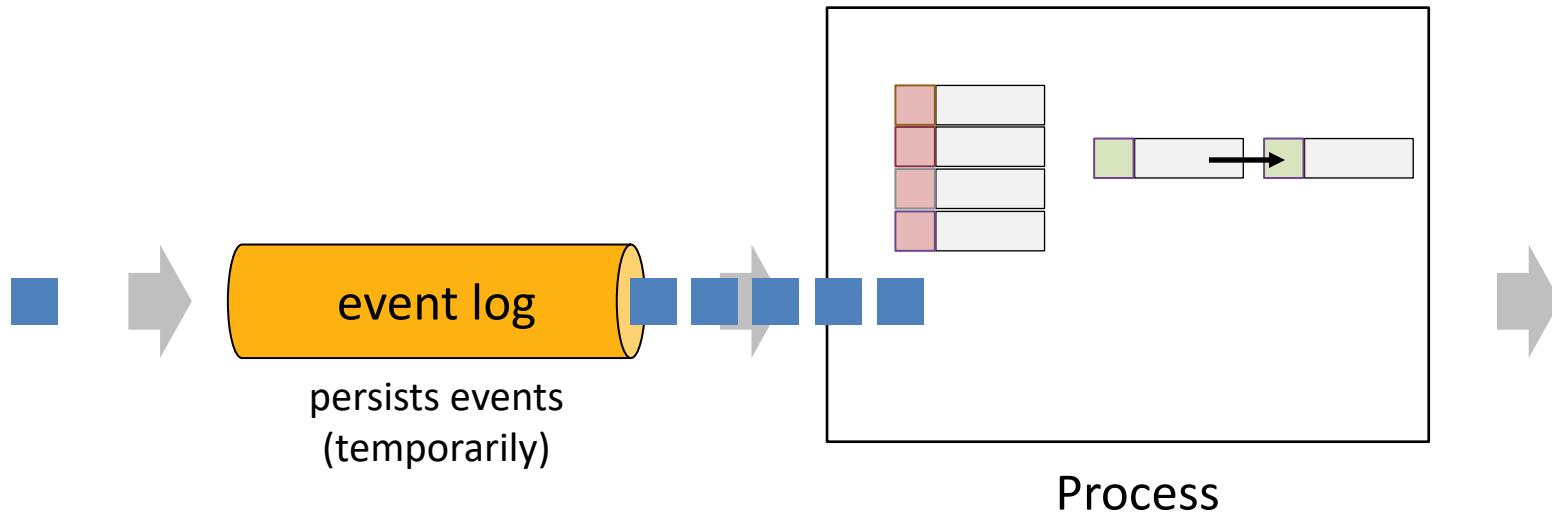
Event Sourcing + Memory Image



Event Sourcing + Memory Image



Recovery: Restore snapshot and replay events
since snapshot



Distributed Memory Image



Distributed application, many memory images.
Snapshots are all consistent together.

